

HOW GTM RADAR IS BUILT

Methodology & Decision Log

The buckets, the filters, the scoring, and every pivotal scope call made during the build.

GTM Radar is a research workflow, not an automated pipeline. Its quality depends on what gets onto the watchlist and what gets thrown out. This page documents both — the operating logic and the decisions that shaped it. Read it as the operator's manual.

Six buckets, one filter

AI is the company filter, not the role filter. A VP Sales hired at a company racing to ship AI is the briefing's target. A 'Head of AI' or 'VP of AI' is not. The six buckets organize the watchlist around where Betts's book actually lives, with a bias toward AI-late incumbents (the largest opportunity) over AI-native cohorts (the smaller, more visible one).

- 01 AI-Native (15 cos)** — AI is core product architecture. Harvey, Writer, Glean, Cresta, Observe.AI, Electric, Gong, Notion, plus seven scaling cohorts in customer support / knowledge work / legal.
- 02 RevTech (12 cos)** — sales / marketing / RevOps tooling. Salesloft, Outreach, Clari, 6sense, Demandbase, Apollo.io, ZoomInfo, Common Room, UserGems, Clay, Sybill, Lavender.
- 03 Fintech + Cyber (10 cos)** — adjacent sectors where AI is reshaping the GTM motion. Ramp, Brex, Mercury, Rho, Abnormal, Vectra, Dropzone, Prophet, Sysdig, Red Points.
- 04 AI-Late Adopter (16 cos, biggest bucket)** — mid-cap SaaS being pushed to adopt AI. Zendesk, DocuSign, Box, Dropbox, Twilio, Intercom, Okta, Atlassian, Asana, Monday, Smartsheet, Zoom, Samsara, Klaviyo, PagerDuty, Procore.
- 05 Incumbent Context (2 cos, context-only)** — HubSpot, Gainsight. Tracked for market awareness; not Betts search targets.
- 06 Pre-Discovery (8 cos)** — vertical-AI companies that match the Betts ICP but sit off the conventional GTM-recruiting beat. Trunk Tools, Mendel.ai, Sixfold, Cleric, 11x, Nooks, Tessel, Speak. The strongest version of this product surfaces deals the firm doesn't know about yet.

Three filters decide what's on the list

All three must be true. The filters are conservative on purpose: they keep the watchlist defensible and the briefing density high.

- 01 Active GTM hiring need.** Series B through pre-IPO, or mid-cap public / PE-backed. No pre-PMF.
- 02 No internal recruiting army or brand magnet.** Anthropic gets 50K inbound applications. The pain has to be real for Betts to win.
- 03 Matches Betts's documented pattern.** Tech / SaaS, GTM-heavy, volume hires or exec search. Anchored in public Betts case studies — Harvey, Zendesk, Grammarly, Sysdig, Tray.io, Electric, Attentive, Staffbase.

Scoring rubric

Every signal is scored 1–5. Default cutoff for inclusion in the briefing is ≥ 3 . The scoring is a forward-looking hiring-likelihood read, not a press-release weight.

SCORE	LABEL	PREDICTS	EXAMPLE
5	Hot	Hires within 30 days	Late-stage raise + AI-native + active reqs
4	Warm	Hires within 60 days	Recent round or new exec at scaling AI co
3	Watch	Hires within 90 days	Product launch or early-stage raise
2	Background	No direct GTM read	Monitor — does not enter the briefing
1	Noise	Drop	Wrong sector, wrong stage, or unverified

Pivotal scope decisions

Every entry is a real decision made during the build. Each changed what gets included or how the briefing reads.

- 01 Cut hyperscalers and brand magnets.** Anthropic, OpenAI, Perplexity, Cursor, Cognition, Poolside, Magic excluded from v1. They don't need Betts. Healthcare AI scoped out — different recruiter network.
- 02 Re-classified Gong and Notion to AI-Native.** Both have AI as core product architecture, not a late-added layer. Moving them out of incumbent-pivot sharpened the bucket distinction.
- 03 Defined the role filter.** AI is the company filter, not the role filter. 'Head of AI' / 'VP of AI' / 'AI Product Manager' explicitly excluded — those aren't Betts's book. A VP Sales at the same company is.
- 04 Added 'too big for Betts' exclusion.** Salesforce and ServiceNow removed. Internal TA armies at scale.
- 05 Created 'context-only' bucket.** HubSpot and Gainsight tracked for market-pitch awareness but not as search targets. Their events appear in TL;DRs / sidebars, not as full signal cards.
- 06 Expanded watchlist by 12 cos after a quiet W17.** First two scans returned zero on-target signals. Added seven AI-native scaling cohorts (Sierra, Decagon, Hebbia, Forethought, Maven AGI, Crescendo, EvenUp), three RevTech (Clay, Sybill, Lavender), two AI-late adopters (PagerDuty, Procore). The expanded scan surfaced four verified launches in the W17 window.

The verification rule

Every signal traces to a primary source — company press release, blog, SEC filing — or a tier-1 secondary (TechCrunch, Reuters, WSJ, Information, Bloomberg) within the published week's window. The brand of this artifact is 'never hallucinates.' One fabricated signal kills the credibility of the briefing.

Rho Series B Page-update on a 2024 announcement	Apollo \$110M Conflated with the 2022 raise	Abnormal FedRAMP 2025 announcement — outside window
Vectra - Phillips CRO Dec 2025 — outside window	Mercury \$5B Rumor — no primary source	Zendesk AI rollout Mar 30 — outside window

6 REJECTED · 5 PUBLISHED · THE REJECTED LIST SHIPS INSIDE EACH BRIEFING.

What this could become

Three directions worth pressure-testing if there's appetite to extend the work. None are built yet. Each is a different shape of bet on what makes recruiter intel actually useful at the desk.

- 01 Two more verticals.** Extend coverage to cohorts Betts already places into but the current sources don't reach — vertical SaaS (Toast, Procore, ServiceTitan-adjacent), developer infra (Vercel, Supabase, Resend), or whichever cohort sits closest to your active searches.
- 02 An interactive query layer.** A recruiter pulls up a target company mid-call and asks 'who's likely available' or 'what's the hiring read here?' Not a dashboard — a conversational query against the

same research substrate these PDFs are built on.

-
- 03 Backtest the rubric.** Take six months of confirmed Betts placements, replay the watchlist as if those placements hadn't happened yet, and check whether the 1–5 scoring would have flagged them in the right window. Turns the 30/60/90-day claim from a hypothesis into a measurement.
-

WHERE I'D LIKE FEEDBACK

Of the three directions above, which one would actually be useful to a Betts recruiter at their desk on a Monday morning? Or — what's missing? The PDFs were the easy part. Building the next version against a real read beats another guess.

Kyle Findlay

kylefindlay828@gmail.com