

HOW GTM RADAR IS BUILT

Methodology & Decision Log

The buckets, the filters, the scoring, and every pivotal scope call made during the build.

GTM Radar is a research workflow, not an automated pipeline. Its quality depends on what gets onto the watchlist and what gets thrown out. This page documents both — the operating logic and the decisions that shaped it. Read it as the operator's manual.

Five buckets, one filter

AI is the company filter, not the role filter. A VP Sales hired at a company racing to ship AI is the briefing's target. A 'Head of AI' or 'VP of AI' is not. The five buckets organize the watchlist around where Betts's book actually lives, with a bias toward AI-late incumbents (the largest opportunity) over AI-native cohorts (the smaller, more visible one).

- 01 AI-Native (15 cos)** — AI is core product architecture. Harvey, Writer, Glean, Cresta, Observe.AI, Electric, Gong, Notion, plus seven scaling cohorts in customer support / knowledge work / legal added April 27.
- 02 RevTech (12 cos)** — sales / marketing / RevOps tooling. Salesloft, Outreach, Clari, 6sense, Demandbase, Apollo.io, ZoomInfo, Common Room, UserGems, Clay, Sybill, Lavender.
- 03 Fintech + Cyber (10 cos)** — adjacent sectors where AI is reshaping the GTM motion. Ramp, Brex, Mercury, Rho, Abnormal, Vectra, Dropzone, Prophet, Sysdig, Red Points.
- 04 AI-Late Adopter (16 cos, biggest bucket)** — mid-cap SaaS being pushed to adopt AI. Zendesk, DocuSign, Box, Dropbox, Twilio, Intercom, Okta, Atlassian, Asana, Monday, Smartsheet, Zoom, Samsara, Klaviyo, PagerDuty, Procore.
- 05 Incumbent Context (2 cos, context-only)** — HubSpot, Gainsight. Tracked for market awareness; not Betts search targets.

Three filters decide what's on the list

All three must be true. The filters are conservative on purpose: they keep the watchlist defensible and the briefing density high.

- 01 Active GTM hiring need.** Series B through pre-IPO, or mid-cap public / PE-backed. No pre-PMF.
- 02 No internal recruiting army or brand magnet.** Anthropic gets 50K inbound applications. The pain has to be real for Betts to win.
- 03 Matches Betts's documented pattern.** Tech / SaaS, GTM-heavy, volume hires or exec search. Anchored in public Betts case studies — Harvey, Zendesk, Grammarly, Sysdig, Tray.io, Electric, Attentive, Staffbase.

Scoring rubric

Every signal is scored 1–5. Default cutoff for inclusion in the briefing is ≥ 3 . The scoring is a forward-looking hiring-likelihood read, not a press-release weight.

SCORE	LABEL	PREDICTS	EXAMPLE
5	Hot	Hires within 30 days	Late-stage raise + AI-native + active reqs
4	Warm	Hires within 60 days	Recent round or new exec at scaling AI co
3	Watch	Hires within 90 days	Product launch or early-stage raise
2	Background	No direct GTM read	Monitor — does not enter the briefing
1	Noise	Drop	Wrong sector, wrong stage, or unverified

Pivotal scope decisions (dated)

Every entry is a real decision made during the build. Each changed what gets included or how the briefing reads.

- 01 **Apr 20 — cut hyperscalers and brand magnets.** Anthropic, OpenAI, Perplexity, Cursor, Cognition, Poolside, Magic excluded from v1. They don't need Betts. Healthcare AI scoped out the same day — different recruiter network.
- 02 **Apr 21 — re-classified Gong and Notion to AI-Native.** Both have AI as core product architecture, not a late-added layer. Moving them out of incumbent-pivot sharpened the bucket distinction.
- 03 **Apr 21 — defined the role filter.** AI is the company filter, not the role filter. 'Head of AI' / 'VP of AI' / 'AI Product Manager' explicitly excluded — those aren't Betts's book. A VP Sales at the same company is.
- 04 **Apr 21 — added 'too big for Betts' exclusion.** Salesforce and ServiceNow removed. Internal TA armies at scale.
- 05 **Apr 22 — created 'context-only' bucket.** HubSpot and Gainsight tracked for market-pitch awareness but not as search targets. Their events appear in TL;DRs / sidebars, not as full signal cards.
- 06 **Apr 27 — expanded watchlist by 12 cos after a quiet W17.** First two scans returned zero on-target signals. Added seven AI-native scaling cohorts (Sierra, Decagon, Hebbia, Forethought, Maven AGI, Crescendo, EvenUp), three RevTech (Clay, Sybill, Lavender), two AI-late adopters (PagerDuty, Procore). The expanded scan surfaced four verified launches in the W17 window.

The verification rule

Every signal traces to a primary source — company press release, blog, SEC filing — or a tier-1 secondary (TechCrunch, Reuters, WSJ, Information, Bloomberg) within the published week's window. The brand of this artifact is 'never hallucinates.' One fabricated signal kills the pitch.

Week 17 published its rejected list as proof of method: Rho Series B (page-update on a 2024 announcement), Apollo \$110M (conflated with the 2022 raise), Abnormal FedRAMP (2025), Vectra · Phillips CRO (Dec 2025), Mercury \$5B (rumor), Zendesk AI agent rollout (Mar 30 — outside window). Six rejected, five published. The rejected list ships inside the briefing.

WHAT TO DO WITH THIS PAGE

Read this before any conversation about scaling GTM Radar inside Betts. The bucket logic, filters, and decision history are the thing that ports — not the PDFs.

Kyle Findlay
kylefindlay828@gmail.com